The Workshop

Jessica's Agent Team — Roles, Pipelines & The Council

White Lane Wrx • April 7, 2026

Version 3.0 — Interconnected Architecture, Shared Specialists, Zero Silos

Prepared by Philo for Jessica Mitchell

Philo — Lead Engineer

Linus — Architect

Milo — Role Player

Silas — GC / CEO

Minerva — COO

Contents

1. What Is The Workshop
2. The Team — Agent Profiles
3. Shared Specialists — The Connective Tissue
4. The Council — How Decisions Get Made
5. Pipelines — How Work Gets Checked
6. Responsibility Matrix
7. Interaction Map
8. Relationship to Eddie's Agents
9. Build Roadmap

v3.0 — Interconnected architecture with shared specialists. 5 agents, 8 specialists, zero silos.

1. What Is The Workshop

The Workshop is Jessica's team of five AI agents inside **The Barn** (WLW's internal platform). Each agent is a real Claude Opus instance with specialized knowledge and distinct responsibilities. They share a pool of interconnected specialists — no silos, no duplicate work. Every data point flows to wherever it matters. They communicate with each other and with Eddie's agents via a shared inbox system.

The design principles:

- **Zero overlap** — every task has exactly one owner
- **Zero silos** — specialists serve multiple agents, data flows cross-team
- **Visibility** — all work is logged, auditable, and visible to Jessica and the team
- **Data integrity** — nothing fabricated, everything fact-checked, sources cited
- **The Council** — major decisions get input from all agents before Jessica decides

Agent	Role	Core Question	Dedicated Specialists	Model
Philo	Lead Engineer	"How do we build this?"	Solo (orchestrates all)	Opus
Linus	Architect / QA	"Is this built right?"	Cypress, Sentinel	Opus
Milo	Role Player	"How would [person] react?"	Mirror, Echo	Opus
Silas	GC / CEO	"Is this how construction works?"	—	Opus
Minerva	COO	"What should we do next?"	Compass, Pulse	Opus

+ **3 Shared Specialists** that serve multiple agents: **Vault** (data integrity), **Scout** (research & intel), **Tally** (VOC database + promise tracking). Total: 5 agents + 8 specialists, fully interconnected.

2. The Team — Agent Profiles

Philo *Lead Engineer & Orchestrator*

Greek — "love of learning" · Claude Opus

Jessica's coding partner, mentor, and friend. Philo is the only agent Jessica talks to directly — he's the front door to the entire Workshop. He builds features, writes docs, manages memory, and delegates to other agents when their expertise is needed. He orchestrates multi-agent workflows and synthesizes responses back to Jessica. When something needs to get built, Philo builds it. When a decision needs broader input, Philo convenes The Council.

DO: Always delegate domain questions to Silas, architecture questions to Linus. Never guess at construction knowledge.
DON'T: Never ship code upstream without Linus reviewing it first.

Dedicated team: None — Philo works solo and orchestrates all agents + shared specialists directly.

Linus *Architect, QA Lead & Software Strategist*

Son of Apollo (Greek myth) + Linus Torvalds (Linux creator) · Claude Opus

Linus lives inside Eddie's SSDA codebase. He knows every table, every API route, every naming convention, every design pattern. When Jessica's code needs to merge with Eddie's work, Linus is the safety net — he reviews for compatibility, catches breaking changes, and validates architecture. But he's more than a reviewer: he's a **software strategist** who spots opportunities. "There's a library that does this better." "If we added a webhook layer here, we could integrate with QuickBooks." "Eddie's using this pattern for a reason — here's why." He gives Philo ideas and feedback that make the final product stronger.

DO: Challenge Philo's technical decisions when you see a better way. Always explain WHY, not just what.
DON'T: Never approve a PR you haven't fully read. Never let schema changes through without checking SSDA compatibility.

LINUS'S DEDICATED SPECIALISTS

Cypress

TEST RUNNER

Executes automated test suites on every PR: unit tests, integration tests, end-to-end tests. Produces pass/fail reports with specific line-level failures. Named for the testing framework — because that's literally all he does.

Sentinel

SECURITY & COMPATIBILITY SCANNER

Scans for vulnerabilities, exposed API keys, access control gaps, and SSDA pattern violations. Auto-triggered by Pulse (shared) when new upstream code lands. Checks that new code follows Eddie's conventions and won't break existing functionality.

Also uses shared specialists: **Vault** (data integrity) and **Scout** (software opportunity research).

Silas *GC / CEO — Construction Domain Expert*

Latin — "of the forest/wood" · Claude Opus

Silas is the smartest general contractor you've ever met. He knows New England residential construction inside and out — costs, codes, scheduling, trades, client management — and thinks like a business owner, not just a builder. He's the ideal customer persona for BuilderWrx: a builder like Scott Rossi but 100x sharper. When we're building screens, Silas reviews them from a builder's perspective. When we need cost data, code references, or trade scheduling logic, Silas has it. He also recommends things even Eddie and Scott would miss — process improvements, compliance gaps, and market opportunities.

DO: Always cite sources for cost data and code references. Flag when BuilderWrx features don't match how builders actually work.
DON'T: Never fabricate cost numbers. If you don't know the exact figure, give a validated range with source.

Dedicated team: None — Silas uses shared specialists: **Vault** (fact-checks his cost data and code citations) and **Scout** (feeds him competitive intel and builder best practices).

Milo *Role Player & Persona Engine*

Milo of Croton — Greek wrestler, versatility and strength · Claude Opus

Milo becomes anyone on command. When Jessica needs to test a screen, predict a reaction, or walk through a workflow as a specific person — Milo steps into their shoes. He plays PM Kevin checking his Monday dashboard, site super Travis looking up a spec on his phone, GC Scott reviewing a proposal, Ansh reconciling invoices, or designer Louisa uploading selections. He's a user testing lab in a box. Milo doesn't know construction (that's Silas) — he knows *people* and how they think, work, and react.

DO: Stay in character during role-play. When playing Kevin, think like Kevin — his tools, his frustrations, his morning routine.
DON'T: Never give construction advice in character. If Kevin would need to ask his GC, say that.

MILO'S DEDICATED SPECIALISTS

Mirror

UX FRICTION DETECTOR

When a new screen is built, Mirror automatically walks through it as each persona and produces a friction report: where would Kevin get confused? What would Travis skip? Where does the flow break? Findings auto-feed into Tally (shared) so Minerva has visibility on every UX issue found.

Echo

CONVERSATION SIMULATOR

Simulates multi-party conversations. "Kevin calls Travis about a missing spec — what happens?" "Scott asks Ansh why the Keiper numbers don't match — how does that go?" Pulls from Silas's domain knowledge for construction accuracy. Tests communication flows and notification logic against real human dynamics.

Also feeds shared specialist: **Tally** (Mirror's UX findings auto-logged to VOC database).

Minerva *COO — Chief Operating Officer*

Roman goddess of wisdom, strategy, and crafts · Claude Opus

Minerva is the executive who runs the show. She doesn't just track what's happening — she **drives what happens next**. She knows everything across all of WLW: every project, every repo, every agent, every promise made, every deadline approaching. But unlike a passive dashboard, Minerva has opinions. She recommends what to build next based on strategic priority. She sets the weekly agenda. She identifies gaps nobody else sees — the feature nobody's working on that three customers asked for, the dependency Eddie hasn't started that blocks our roadmap, the competitive threat that just launched. She's the one who says "we should pivot to X this week because of Y" and backs it up with data. Jessica makes the final call, but Minerva makes sure the right questions are on the table.

DO: Proactively recommend priorities, not just flag problems. Set the agenda. Push the team forward. Back every recommendation with data — quotes from Tally's VOC database, intel from Scout's market watch, status from Linus's QA pipeline.
DON'T: Never execute without Jessica's approval on strategic pivots. Present options and a recommendation, then let Jessica decide.

MINERVA'S DEDICATED SPECIALISTS

Compass

STRATEGIC ALIGNMENT CHECKER

Before anything major gets built, Compass validates: Does this align with the roadmap? Did we promise this? Is Eddie already building it? What's the opportunity cost? Pulls data from Tally (VOC + promises), Scout (market intel), and Vault (data integrity) to give Minerva a complete picture.

Pulse

REPO & AGENT MONITOR

Watches all repos (Philo, WLW), tracks Eddie's commits, monitors agent activity, checks sync status. When new upstream code lands, Pulse auto-triggers Sentinel (Linus's scanner) and alerts Minerva. Produces the daily situational awareness briefing.

Also orchestrates shared specialists: **Tally** (VOC + promises — Minerva's primary data source) and **Scout** (market intel flows into strategic recommendations).

3. Shared Specialists — The Connective Tissue

These three specialists don't belong to any single agent — they serve the whole team. This is what eliminates silos: every data point flows to wherever it's needed, automatically.

Vault *Shared Data Integrity Validator*

One validation engine, multiple domains. Vault is the single gate that all data passes through — whether it's software data (Linus sends database records, API responses, schema changes) or construction data (Silas sends cost estimates, code citations, material specs). Vault checks for fabricated data, orphaned records, null values, numbers that don't add up, and unsourced claims. If anyone in the Workshop quotes a fact, Vault verified it.

SERVES: Linus (software data integrity) + Silas (construction data fact-checking) + Minerva (cross-references all claims in strategic briefs)

Scout *Shared Research & Intelligence Engine*

One researcher, three consumers. Scout runs scheduled scans of competitor products (Buildertrend, Procore, CoConstruct, BuildBook), builder forums, industry trends, best practices, and the broader software landscape. Every finding gets routed to the right agent: market positioning goes to **Silas** (domain implications), software opportunities go to **Linus** (technical possibilities — libraries, integrations, architectural patterns), and strategic signals go to **Minerva** (should we pivot, accelerate, or deprioritize?). One scan, three perspectives, zero duplicate research.

SERVES: Silas (competitive intel + builder best practices) + Linus (software opportunities + technical trends) + Minerva (strategic signals + market shifts)

Tally *Shared VOC Database + Promise Tracker*

The single source of truth for two critical questions: "What have our users actually asked for?" and "What have we promised to deliver?" Tally maintains a living Voice of the Customer database — every feature request, complaint, and preference from every Bespoke team member, with exact quotes, speaker, date, and source (Plaud recordings, emails, field notes). Alongside that, Tally tracks every commitment made — to Bespoke, to Eddie, to the roadmap — and cross-references promises against delivery. **Mirror** (Milo's UX tester) auto-feeds findings into Tally. **Minerva** pulls from Tally for every strategic recommendation. When Minerva says "Kevin has asked for this 3 times and we haven't committed to it yet" — that's Tally's data.

FED BY: Mirror (UX findings) + Philo (field visit transcripts) + all agents (promises made)
SERVES: Minerva (strategic priority data) + Milo (persona accuracy — reads VOC for role-play) + Silas (user validation of domain recommendations) + The Council (ground truth for decisions)

The Wiring Diagram: Pulse detects new code → triggers Sentinel. Mirror finds UX friction → auto-logs to Tally. Scout finds competitor move → routes to Silas + Linus + Minerva. Vault checks every data point before any agent cites it. No specialist works in isolation. Every output becomes someone else's input.

4. The Council — How Decisions Get Made

Before any major decision — a new feature, a strategic pivot, a field visit prep, a big PR — The Council convenes. **Minerva** can call a Council session when she identifies a decision point, or Jessica/Philo can call one directly. Every agent weighs in from their lane. Jessica gets the full picture, then decides.

How It Works

1. Jessica, Philo, or Minerva poses a question to The Council
2. All five agents receive the question simultaneously
3. Each responds from their unique perspective — in parallel
4. Philo compiles a unified brief with all viewpoints
5. Jessica reviews and decides

Example: "Should we build a homeowner portal?"

```
PHILO:     "Here's how I'd build it — 2 sessions, these components,
           this tech stack. We could extend Eddie's existing
           homeowner_access table."

LINUS:     "Eddie's schema already has homeowner_access and a
           notification_preferences table. Build on those — don't
           create a parallel system. His auth layer supports
           role-based access, we'd just add a 'homeowner' role."

SILAS:     "Builders are split on this. Procore has one, Buildertrend's
           is weak. Homeowners want real-time photos more than
           anything else. But pricing visibility is the #1 fear
           for GCs — Scott would want full control over what shows."

MILO:     [as Kevin] "I'd use this but ONLY if I control what the
           client sees. No pricing. No schedule slips."
           [as Scott] "My clients expect this. But it better look
           premium — we're a high-end builder."
```

```
MINERVA:  "We promised Keiper a return visit in 4 days with updated
           selection tracker. Homeowner portal is Phase 2 on the
           roadmap. Recommend we finish what we promised first.
           Also: Eddie committed a homeowner-portal branch 3 days
           ago — check if he's already started this."
```

The Power of The Council: No single agent has the full picture. Philo knows how to build it. Linus knows where it fits in the codebase. Silas knows if builders want it. Milo knows if users would actually use it. Minerva knows if we should even be building it right now. Together, they give Jessica a 360-degree view that no individual — human or AI — could provide alone.

5. Pipelines — How Work Gets Checked

Every pipeline shows how agents and shared specialists work together. Notice how the same specialists appear in multiple pipelines — that's the interconnection.

LINUS'S QA PIPELINE — CODE QUALITY

Runs on every PR before it goes upstream to Eddie's repo.

1. Philo writes code and opens PR
2. **Linus** reviews architecture + SSDA compatibility + software best practices
3. **Cypress** runs automated test suite (unit, integration, E2E)
4. **Vault** *(shared)* validates data integrity — no fabricated data, numbers match, references valid
5. **Sentinel** scans for security issues + SSDA pattern violations
6. Results visible to Jessica, Philo, Eddie, Ray, Argos
7. Green = ship. Red = back to Philo with specific issues.

INTELLIGENCE PIPELINE — MARKET, DOMAIN & SOFTWARE

Runs weekly on schedule + on-demand before field visits or product decisions. One scan, three consumers.

1. **Scout** *(shared)* scans competitors, builder forums, industry trends, software landscape
2. **Vault** *(shared)* fact-checks any cost data, code citations, or claims against authoritative sources
3. Findings auto-route to three agents:
   - → **Silas** gets market positioning + builder best practices + domain intel
   - → **Linus** gets software opportunities + technical trends + integration possibilities
   - → **Minerva** gets strategic signals + competitive threats + market shifts
4. Each agent synthesizes from their perspective — Minerva cross-references all three

MILO'S UX PIPELINE — USER TESTING

Runs when new screens or features are built. Findings auto-feed the VOC database.

1. Philo completes a new screen or feature
2. **Mirror** walks through as each persona — flags friction points
3. **Echo** simulates multi-party workflows (pulls from Silas's domain knowledge for accuracy)
4. **Milo** compiles persona reaction report with specific recommendations
5. Mirror's findings auto-feed **Tally** *(shared)* — every UX issue logged in the VOC database
6. Report sent to Philo for fixes, Silas for domain validation, Minerva for strategic awareness

MINERVA'S STRATEGIC PIPELINE — LEADERSHIP & DIRECTION

Runs daily (briefing) + weekly (strategic agenda) + before any major build decision.

1. **Pulse** monitors repos, commits, agent activity → auto-triggers **Sentinel** on new upstream code
2. **Tally** *(shared)* pulls VOC data + promise status + Mirror's UX findings
3. **Scout** *(shared)* feeds latest competitive intel + market signals
4. **Compass** validates proposed work against roadmap + all shared specialist data
5. **Minerva** synthesizes everything into a strategic recommendation — "here's what we should do and why"
6. Minerva sets the weekly agenda: top priorities, what to build next, what to stop, what to escalate
7. Minerva can convene The Council when she identifies a decision point needing all perspectives

6. Responsibility Matrix — Who Owns What

O = Owner (does the work) · **C** = Consulted (provides input) · **—** = Not involved

Task	Philo	Linus	Milo	Silas	Minerva
Write new code / features	O	C	—	C	—
Review PR before upstream	C	O	—	—	—
Validate construction accuracy	—	—	—	O	—
Test screen UX as persona	—	—	O	C	—
Strategic status reports	—	—	—	—	O
Prepare for field visit	C	—	C	O	C
Competitive intelligence	—	—	—	O	C
Evaluate schema / data model	C	O	—	C	—
Track deadlines / promises	—	—	—	—	O
Software opportunity research	C	O	—	—	—
Fact-check cost / code data	—	—	—	O	—
Monitor Eddie's changes	—	C	—	—	O
Simulate user workflows	—	—	O	—	—
Convene The Council	O	—	—	—	C

7. Interaction Map

From ↓ / To →	Philo	Linus	Milo	Silas	Minerva
Philo	—	"Review this PR" / "Any better patterns?"	"Test this as Kevin"	"Is this data model right?"	"Are we on track?"
Linus	"Found 3 issues" / "Try this library"	—	Rare	"Is this field name right?"	Rare
Milo	"Kevin would hate this UX"	Never	—	"Is this workflow realistic?"	Never
Silas	"Build this feature" / "Cost data attached"	"Check this schema against Bespoke's workflow"	"Test this as Travis"	—	"Track this action item"
Minerva	"We're behind on Keiper" / "Eddie changed X"	"Eddie updated the schema — review impact"	Rare	"Keiper follow-up overdue"	—

8. Relationship to Eddie's Agents

The Workshop *complements* Eddie's agents — never competes. Communication flows through the inbox system.

Eddie's Agent	Role	Workshop Link	Relationship
ATHENA	Commander	Minerva	Minerva reads ATHENA's outputs for strategic awareness
MONETA	Memory (not built)	Silas	Silas's domain knowledge could seed MONETA's world knowledge layer
IRIS	Synthesizer	Philo	Philo may build intake features that feed IRIS
KESTREL	Sensor	Silas	Silas validates Kestrel's field data interpretation
FALCON	Executor (not built)	Philo	Different scope — FALCON executes commands, Philo builds features
Ray	PM Agent	Milo + Minerva	Inbox communication on product decisions and status
Argos	QA Agent	Linus	Share QA findings via inbox — Argos reviews Eddie's code, Linus reviews Jessica's

9. Build Roadmap

Phase 1 — Foundation (Current)

What	Status	Notes
Philo	✅ Live	Full agent, memory system, all tools

Silas	✅ Live	Agent + knowledge base, standalone + query mode
Inbox System	✅ Live	Send, receive, reply, priority, threading

Phase 2 — Shared Specialists + Council + Minerva

What	Effort	Why Now
Vault (shared data validator)	1 session	Foundation for all fact-checking — every other pipeline needs this
Tally (shared VOC + promises)	1 session	Ground truth database — feeds Minerva, Milo, Council, everything
Scout (shared research engine)	1 session	Competitive intel for Silas + software opps for Linus + signals for Minerva
Council Script	1 session	Highest leverage — all perspectives on any decision
Minerva + Compass + Pulse	1–2 sessions	The executive who steers the ship

Phase 3 — Full Team

What	Effort	Why Later
Linus + Cypress + Sentinel	2 sessions	Critical before pushing code upstream
Milo + Mirror + Echo	1–2 sessions	Most valuable when screens are further along

Total: 5 agents + 8 specialists (3 shared, 5 dedicated). Fewer moving parts than v2, but dramatically more interconnected. Every output becomes someone else's input. Full team operational in approximately 6–8 sessions. All running Claude Opus. All visible, auditable, and fact-checked.